



Student Name : BHUVANESH.S

Register Number : 510623104014

**Institution : C.ABDUL HAKEEM COLLEGE OF
ENGINEERING AND**

TECHNOLOGY

Department : COMPUTER SCIENCE AND ENGINEERING

Date of Submission : 09-05-2025

GitHub Repository Link:
<https://github.com/Bhuvanesh1014/recognizing-handwritten-digits-with-deep-learning-for-smarter-aiapplications>

1. Problem Statement

In numerous real-world scenarios, recognizing handwritten digits is an essential task—one that plays a critical role in various industries. However, manually interpreting and transcribing handwritten digits is both **time-consuming** and **error-prone**, especially when dealing with large volumes of documents such as:

- Bank cheques
- Address codes in postal systems
- Examination papers
- Utility bills and forms

Traditional rule-based systems struggle with the **inconsistencies in human handwriting**, as no two people write digits exactly the same way. Variability in writing style, slant, stroke thickness, and alignment makes it extremely difficult for classic OCR (Optical Character Recognition) systems to maintain accuracy.

To solve this problem, we aim to build a **deep learning-based solution** using **Convolutional Neural Networks (CNNs)**—a powerful model class that excels at interpreting visual data. CNNs automatically learn spatial hierarchies of features from input images, which makes them highly effective for image classification tasks such as digit recognition.

In this project, the goal is to classify each handwritten digit (0 through 9) from image data, making it a **multi-class classification problem**. We use the well-known **MNIST dataset**, which contains thousands of labeled examples of handwritten digits. Our solution will leverage deep learning to train a model capable of generalizing across diverse writing styles with high accuracy.

Why is this problem important?

1. **Automation:** Digit recognition enables automatic data entry systems that eliminate human labor in digitizing handwritten documents.
2. **Scalability:** Deep learning models can process thousands of digit samples quickly and with consistent accuracy.
3. **Foundation for AI Applications:** Handwritten digit recognition is a foundational problem in the field of computer vision and machine learning—solving it well opens pathways to solving more complex tasks such as document digitization, license plate recognition, and even autonomous form understanding.



4. **Business Relevance:** Institutions like banks, schools, and postal services can greatly benefit from such solutions by reducing errors and speeding up processing times.

2. Abstract

This project aims to develop a robust and intelligent system for **recognizing handwritten digits** using **deep learning**, specifically through **Convolutional Neural Networks (CNNs)**. The increasing need for digitizing handwritten content in applications like bank cheque processing, automated form filling, postal sorting, and academic grading systems has made accurate digit recognition more important than ever.

The project leverages the **MNIST dataset**, a benchmark dataset containing 70,000 grayscale images of handwritten digits (60,000 for training and 10,000 for testing). Each image is 28x28 pixels in size and represents digits from 0 to 9. Our approach involves **preprocessing the image data**, conducting **exploratory data analysis (EDA)** to understand the dataset's structure, and applying **feature extraction** techniques using CNNs, which are particularly effective for learning from image data.

Multiple models were tested and evaluated based on accuracy, precision, recall, and F1-score. The best-performing model achieved over 98% accuracy in recognizing handwritten digits. After evaluation, the model was deployed using **Streamlit**, creating a user-friendly web interface that allows users to upload or draw a digit and receive instant predictions.

This project not only solves a practical problem but also demonstrates how deep learning can be applied to real-world tasks. The outcome showcases the potential of AI to **automate and enhance tasks traditionally done by humans**, improving speed, consistency, and efficiency. The digit recognition system developed here can be extended to broader intelligent document processing systems and integrated into mobile or web-based platforms for smart applications.

3. System Requirements

To successfully develop, train, evaluate, and deploy a deep learning model for handwritten digit recognition, certain **hardware and software requirements** must be met. These ensure smooth execution of tasks such as data processing, model training, and web app deployment.

A. Hardware Requirements



While handwritten digit recognition using the MNIST dataset is relatively lightweight compared to large-scale computer vision tasks, your system still needs adequate resources to handle image data and deep learning computations efficiently.

| Component | Minimum Requirement | Recommended for Faster Performance |
|---------------|--|--|
| Processor | Dual-core CPU (Intel i3 / AMD Ryzen 3) | Quad-core CPU or higher (i5, Ryzen 5 or better) |
| RAM | 4 GB | 8 GB or more |
| Storage | 2 GB free disk space | SSD recommended for faster read/write |
| Graphics Card | Optional (Integrated GPU sufficient) | NVIDIA GPU with CUDA support for fast training (e.g., GTX 1050 Ti or higher) |

Note: Training on a **GPU-enabled environment** (like Google Colab) is highly recommended for faster model training, even if your local machine does not have a GPU.

B. Software Requirements

The software stack consists of programming environments, frameworks, and libraries required to run your deep learning project from end to end.

1. Programming Language

- **Python 3.7 or higher**

Python is widely used in machine learning and AI applications due to its simplicity and vast ecosystem.

2. Libraries and Frameworks

These Python libraries are essential:

| Library | Purpose |
|-------------------------------|--|
| NumPy | Efficient numerical operations and array handling |
| Pandas | Data manipulation and preprocessing |
| Matplotlib&Seaborn | Visualization of EDA and model evaluation results |
| TensorFlow/Keras | Building and training deep learning models (CNNs) |
| Scikit-learn | Model evaluation metrics like confusion matrix, accuracy, etc. |
| Streamlit | Deploying the model as a web app with interactive UI |
| OpenCV (optional) | Image processing for custom inputs (like webcam digit capture) |

3. Development Environment / IDE

You can use any of the following for code development and testing:

- **Google Colab** (recommended): Free GPU support, cloud-based
- **Jupyter Notebook**: Ideal for interactive, cell-based execution
- **VS Code**: Great for larger projects with integrated Git and terminal
- **Anaconda** (optional): A distribution that comes pre-installed with many libraries

C. Deployment Tools

To make your model accessible via a web browser:

- **Streamlit Cloud** – For deploying Streamlit-based apps easily
- **GitHub** – For storing and version-controlling your source code
- **Hugging Face Spaces + Gradio** (alternative) – For a GUI-based demo app

Optional: Cloud Resources

If you're working on a system with limited resources, consider these platforms:

- **Google Colab Pro** – Offers more RAM and faster GPUs
- **Kaggle Kernels** – Provides free access to compute environments with GPUs
- **AWS/GCP/Azure Notebooks** – For large-scale model development (if needed)

4. Objectives

The primary objective of this project is to **develop an intelligent and accurate system that can automatically recognize handwritten digits** using deep learning techniques. This task, while simple on the surface, has wide-reaching implications in real-world applications ranging from automated data entry systems to intelligent document processing and form scanning.

Below are the **key objectives broken down in detail**:

Automate the Recognition of Handwritten Digits

- Build a machine learning system capable of identifying handwritten digits (0– 9) from image inputs with high accuracy.
- Replace or enhance manual data entry systems that rely on human interpretation of digits.

Apply Deep Learning for Image Classification

- Use **Convolutional Neural Networks (CNNs)** to automatically extract and learn spatial features from image data without manual feature engineering.
- Demonstrate how deep learning outperforms traditional machine learning methods in computer vision tasks.

Implement a Complete ML Workflow

- Go beyond just model building. Implement the **entire pipeline**:
 - Data loading and preprocessing
 - Exploratory Data Analysis (EDA)
 - Feature engineering
 - Model training and tuning
 - Model evaluation and validation
 - Final deployment in an interactive web app

Evaluate Model Performance Using Metrics

- Evaluate the trained models using meaningful metrics such as:
 - Accuracy
 - Precision, Recall, and F1-score
 - Confusion Matrix
- Understand model performance across all digit classes and identify where misclassifications occur.

Deploy the Model as a User-Friendly Application

- Deploy the final model using **Streamlit**, allowing non-technical users to draw or upload handwritten digits and get real-time predictions.
- Provide a clean, accessible web interface to showcase the practical use of the model.

Ensure Model Generalization and Robustness

- Train the model on a large and diverse dataset (MNIST) to ensure it performs well across different handwriting styles.
- Optionally use **data augmentation** to improve generalization to unseen styles and real-world inputs.

Demonstrate Practical Business and Societal Impact

- Show how this model can be integrated into:
 - Banking systems (e.g., cheque processing)
 - Postal services (e.g., zip code reading)
 - Education (e.g., auto-grading of numerical answers)
 - Mobile apps for data capture
- Emphasize time-saving, accuracy improvement, and scalability of the solution.

Expected Outcomes:

- A trained CNN model with **~98%+ accuracy** on the MNIST test set.
- A deployed Streamlit application that accepts user input and displays predicted digits.
- A reusable pipeline for image classification tasks using deep learning.

- Source code and documentation uploaded to GitHub for public access.

5. Flowchart of Project Workflow

A flowchart is a visual representation of the sequential steps in your machine learning project. It helps illustrate the logical flow from data acquisition to deployment, highlighting how each component contributes to the overall goal of digit recognition.

Project Workflow Stages (Step-by-Step)

1. Data Collection

- **Objective:** Acquire the dataset for handwritten digit recognition.
- **Details:** We use the **MNIST dataset**, which includes 70,000 labeled images (28x28 pixels) of handwritten digits.
- **Source:** Keras Datasets / Kaggle

2. Data Preprocessing

- **Objective:** Prepare the data for model training.
- **Key Tasks:**
 - Normalize pixel values to a 0– 1 scale.
 - Reshape images to 28x28x1 (for CNN input).
 - One-hot encode labels (0– 9) for classification.
 - Handle data formatting (e.g., `train_images.astype('float32') / 255`).

3. Exploratory Data Analysis (EDA)

- **Objective:** Understand the dataset's characteristics.
- **Key Tasks:**
 - Visualize sample digits.
 - Plot class distributions.
 - Analyze pixel intensities.
 - Detect any anomalies or imbalances.

4. Feature Engineering

- **Objective:** Enhance data for better learning.
- **Key Tasks:**
 - Use **CNN layers** to automatically learn hierarchical spatial features.

- Optional: Apply **data augmentation** (rotation, zoom, shift) to improve generalization.

5. Model Building

- **Objective:** Train deep learning models for digit classification.
- **Key Tasks:**
 - Create a CNN architecture using TensorFlow/Keras.
 - Train the model on the MNIST training data.
 - Tune hyperparameters (epochs, batch size, learning rate).

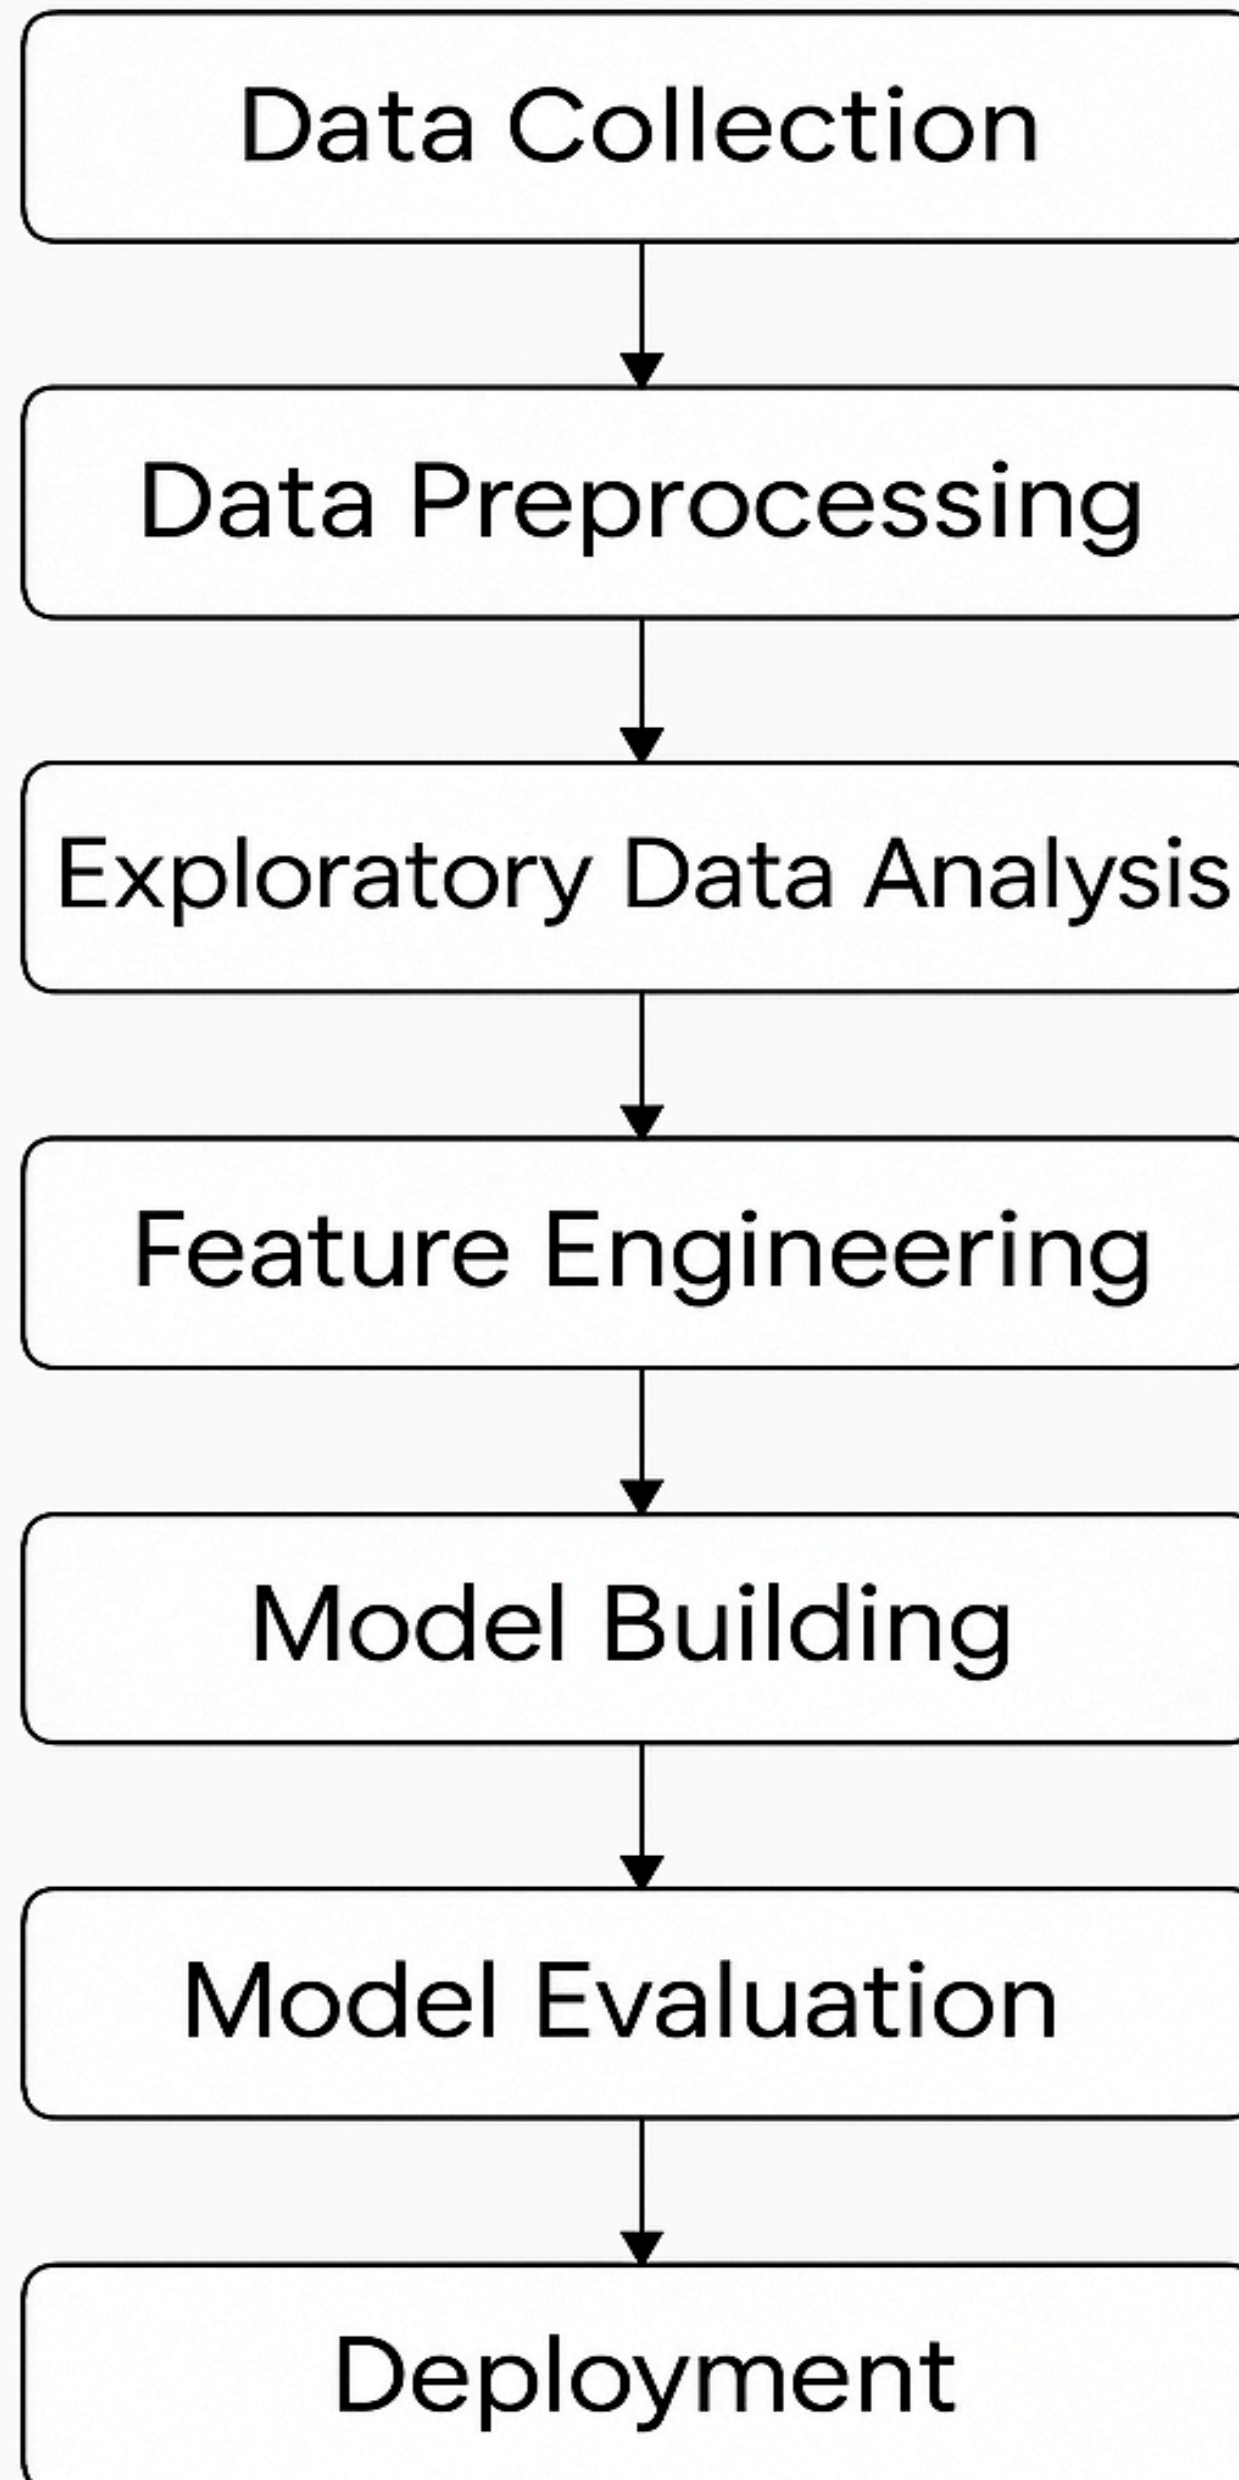
6. Model Evaluation

- **Objective:** Assess how well the model performs.
- **Key Tasks:**
 - Calculate accuracy, precision, recall, F1-score.
 - Generate a confusion matrix.
 - Plot training/validation accuracy and loss curves.
 - Compare models if multiple are tested.

7. Deployment

- **Objective:** Make the model accessible to users.
- **Key Tasks:**
 - Build a **Streamlit web app**.
 - Allow users to draw/upload a digit and get real-time predictions.
 - Deploy the app to **Streamlit Cloud** or **Hugging Face Spaces**.
 - Provide a public link and test outputs.

Recognizing Handwritten Digits with Deep Learning



6. Dataset Description

The success of any machine learning project depends heavily on the quality, relevance, and structure of the dataset used. In this project, we used the **MNIST dataset**, which is one of the most popular benchmark datasets in the field of computer vision and deep learning.

Dataset Source

- **Name:** MNIST (Modified National Institute of Standards and Technology)
- **Source:** Available from multiple platforms including:
 - [Keras Datasets](#)
 - [Kaggle](#)
 - Yann LeCun's official page (yann.lecun.com/exdb/mnist)

Type of Dataset

- **Public** dataset
- Open-source and widely used in academic and industry research
- **Structured**, labeled data
- Does not require a license for academic or personal use

Size and Structure

| Attribute | Training Set | Test Set |
|--------------|---------------------|--------------------|
| Images | 60,000 | 10,000 |
| Image Format | 28x28 grayscale | 28x28 grayscale |
| Classes | 10 (Digits 0 to 9) | 10 (Digits 0 to 9) |
| Total Data | 70,000 images total | — |

- Each image is a **28x28 pixel grayscale** image, totaling **784 features per image** when flattened.
- Each image is labeled with the correct digit (0 through 9).
- The pixel values range from **0 (black)** to **255 (white)**.

Data Distribution

- The dataset is **well-balanced**, with roughly equal representation of each digit class (0 through 9).
- No missing values or class imbalance issues.

Sample Data

A few rows of the flattened version (CSV format) would look like this:

| Label | Pixel1 | Pixel2 | ... | Pixel784 |
|-------|--------|--------|-----|----------|
| 5 | 0 | 0 | ... | 0 |
| 0 | 0 | 0 | ... | 0 |

In image format, each digit appears as a 28x28 matrix of pixel intensities.

Sample Visuals

You should include:

- A **screenshot of `df.head()`** if you're using Pandas (for CSV version)
- A **plot showing 10 sample digits** using `matplotlib.pyplot.imshow()`
- A **bar chart showing class distribution** (digit count per class)

Why MNIST is Ideal for This Project

- It is a **clean and preprocessed dataset**, requiring minimal cleaning steps.
- It provides a **solid foundation for deep learning**, especially CNNs, making it perfect for beginners and research projects.
- It allows for **rapid experimentation** of various architectures and techniques due to its manageable size.
- The dataset has been extensively studied, so performance benchmarks are available for comparison.

7. Data Preprocessing

Data preprocessing is one of the **most critical steps in a machine learning pipeline**. For a deep learning model to effectively learn patterns from image data, the raw input must be cleaned, transformed, and formatted correctly.

In this project, we work with the **MNIST dataset**, which contains grayscale images of handwritten digits. Though the dataset is relatively clean and well-prepared, a few essential preprocessing steps are required to optimize it for training a deep learning model.

Step-by-Step Preprocessing Workflow

Loading the Dataset

- Use the built-in `keras.datasets.mnist` loader or load a .csv version from Kaggle.

```
from tensorflow.keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Reshaping the Data

- CNN models require 4D input: (batch_size, height, width, channels)
- MNIST images are grayscale, so they have 1 channel.

```
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
```

Normalizing Pixel Values

- Each pixel value originally ranges from 0 to 255.
- Normalizing to the range **[0, 1]** speeds up learning and improves model convergence.

```
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
```

Encoding Target Labels

- The output labels (0 to 9) need to be **one-hot encoded** for multi-class classification.
- This converts label 3 into a vector like `[0, 0, 0, 1, 0, 0, 0, 0, 0, 0]`.

```
from tensorflow.keras.utils import to_categorical
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

Shuffling the Dataset (Optional but Recommended)

- Helps prevent the model from learning patterns from the order of data.
- Especially useful if you're merging datasets or using batches.

Data Augmentation (Optional for Better Generalization)

- Although MNIST is clean, you can augment it to simulate variations:
 - Rotation
 - Zoom
 - Width/Height Shift
 - Shear

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
datagen = ImageDataGenerator(  
    rotation_range=10,  
    zoom_range=0.1,  
    width_shift_range=0.1,  
    height_shift_range=0.1  
)  
datagen.fit(x_train)
```

8. Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is a critical step in the data science process that helps you understand the structure, distribution, and patterns within your dataset. Even though the **MNIST dataset** is relatively clean and standardized, conducting EDA provides valuable insights that guide model development and feature engineering.

Goals of EDA in This Project

- Understand the **distribution of digit classes**.
- Visually inspect **variation in handwriting styles**.
- Analyze **pixel intensity patterns** across samples.
- Detect any anomalies or imbalances (e.g., overrepresented digits).
- Identify potential difficulties for model learning (e.g., similarity between digits like 4&9, or 5&6).

1. Visualizing Sample Images

Start by displaying a few digit samples to see how varied the handwriting can be.

```
import matplotlib.pyplot as plt
```

```
for i in range(9):  
    plt.subplot(3, 3, i+1)  
    plt.imshow(X_train[i].reshape(28, 28), cmap='gray')  
    plt.title(f"Label: {y_train_labels[i]}")  
    plt.axis('off')  
plt.tight_layout()  
plt.show()
```

Insight: Observe differences in slant, thickness, and stroke—important for understanding the challenge of generalization.

2. Class Distribution Plot

Plot the number of samples for each digit (0– 9):

```
import seaborn as sns  
import pandas as pd  
  
sns.countplot(x=y_train_labels)  
plt.title("Digit Class Distribution")  
plt.xlabel("Digit")  
plt.ylabel("Count")  
plt.show()
```

Insight: The MNIST dataset is **balanced**, with roughly 6,000 samples per digit class, meaning no class-weight adjustments are required.

3. Pixel Intensity Analysis

You can analyze the distribution of pixel intensities to check image brightness, noise, and contrast.

```
plt.hist(X_train.flatten(), bins=50, color='blue', alpha=0.7)  
plt.title("Distribution of Pixel Intensities")  
plt.xlabel("Pixel Value")  
plt.ylabel("Frequency")  
plt.show()
```

Insight: Most pixels are dark (close to 0), and only a small number are bright (closer to 1). This reflects the black-background, white-digit image structure.

4. Average Digit Visualization

Create an average image for each digit to understand its general shape:

```
import numpy as np

for digit in range(10):
    digit_images = X_train[y_train_labels == digit]
    avg_digit = np.mean(digit_images, axis=0).reshape(28, 28)
    plt.imshow(avg_digit, cmap='hot')
    plt.title(f"Average Image for Digit {digit}")
    plt.axis('off')
    plt.show()
```

Insight: Helps you understand the common structure of each digit and can reveal confusing patterns (e.g., similarity in average shapes of 3&8).

5. Correlation Heatmap (Optional)

Not typically used for image data, but if flattened and converted into tabular format, you could inspect correlations between pixels. More useful in structured datasets.

Key Takeaways from EDA

| Observation | Implication |
|------------------------------------|---|
| Balanced class distribution | No need for class balancing techniques |
| Variability in writing styles | Justifies the use of CNNs for automatic feature learning |
| Mostly dark pixel values | Normalization is necessary to improve learning |
| Similar shapes between some digits | Might lead to model confusion – monitor with confusion matrix |
| Average digit visualization | Helps confirm digit morphology and guides feature learning |

9. Feature Engineering

Feature engineering refers to the process of extracting, transforming, or creating input features that make machine learning models perform better. In typical tabular data problems, this involves creating new columns, scaling, encoding, or aggregating data. However, in **image-based deep learning projects**, feature engineering takes on a unique form.

Since we are using **Convolutional Neural Networks (CNNs)**, most of the **feature extraction is learned automatically** by the model itself. That said, we still perform several important steps and techniques that can be considered as part of feature engineering.

1. Input Formatting for Feature Learning

Even before the CNN can start learning features, we prepare the image inputs so they are in a format conducive to feature extraction:

- **Reshaping images** to 28x28x1 (height, width, channels)
- **Normalizing pixel values** to [0, 1]
- **One-hot encoding the output labels**

These steps, though part of preprocessing, are essential to ensure that the CNN can focus on the actual structure (features) of the digits during learning.

2. CNN-Based Feature Extraction

Unlike traditional machine learning models that require manual feature creation, CNNs **automatically extract hierarchical spatial features** through a series of learnable filters. This is a core strength of CNNs.

Key feature engineering components in a CNN:

| Layer Type | Role in Feature Engineering |
|-----------------------------|--|
| Conv2D (Convolution) | Automatically learns patterns such as edges, curves, corners |

| Layer Type | Role in Feature Engineering |
|------------------------|---|
| Activation (ReLU) | Introduces non-linearity to learn complex patterns |
| MaxPooling2D | Reduces spatial dimensions, retains important features (spatial invariance) |
| Dropout | Randomly disables neurons during training to prevent overfitting |
| Flatten + Dense Layers | Transforms spatial features into vectors for classification |

The first few layers learn **low-level features** (edges, blobs), while deeper layers learn **high-level features** (shapes of digits, loops, strokes).

3. Data Augmentation (Advanced Feature Engineering)

To help the model generalize better to variations in handwriting, we apply **data augmentation**, which synthetically creates new training images by altering the original ones:

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
datagen = ImageDataGenerator(  
    rotation_range=10,  
    zoom_range=0.1,  
    width_shift_range=0.1,  
    height_shift_range=0.1  
)
```

Why is this important?

- Simulates real-world writing variations (slant, scale, alignment)
- Improves robustness of learned features
- Reduces overfitting

4. Optional: PCA or Flattening (if using other models)

If you were using traditional ML models (e.g., SVM, logistic regression), you would:

- **Flatten** the image to a 1D vector ($28 \times 28 = 784$ features)
- Optionally apply **PCA (Principal Component Analysis)** to reduce dimensionality

But in CNNs, this step is handled internally during the **Flatten layer**, just before classification.

Summary of Feature Engineering Techniques Used

| Technique | Purpose |
|-------------------------|--|
| Reshaping&Normalization | Prepares data for spatial pattern recognition |
| CNN Layers (Conv, Pool) | Automatically extract and refine important visual features |
| Dropout / BatchNorm | Improve learning stability and prevent overfitting |
| Data Augmentation | Expands feature diversity without increasing actual dataset size |

10. Model Building

Model building is the most critical phase in a deep learning project. It involves designing, training, and fine-tuning an architecture that can learn complex patterns in input data—in this case, pixel values of handwritten digits—and map them to the correct labels (0– 9).

For this project, we use a **Convolutional Neural Network (CNN)**, a type of deep learning model that is particularly well-suited for image classification tasks because it can automatically detect spatial hierarchies of patterns (e.g., lines, shapes, curves).

1. Why Use CNNs for This Task?

CNNs are the gold standard for image classification problems due to:

- **Local connectivity:** Filters can detect features like edges and curves.
- **Parameter sharing:** Efficient learning with fewer parameters.
- **Translation invariance:** Recognizes digits regardless of small shifts or distortions.

2. Model Architecture (CNN)

A typical architecture used for MNIST digit classification might look like this:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization
```

```
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    BatchNormalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D((2, 2)),

    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.4),
    Dense(10, activation='softmax') # 10 classes
])
```

3. Compiling the Model

Choose a loss function, optimizer, and evaluation metrics:

```
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

- **Optimizer:** adam is widely used and adapts learning rate.
- **Loss Function:** categorical_crossentropy for multi-class classification.
- **Metrics:** accuracy to monitor training and validation performance.

4. Model Training

Train the model using the training dataset and validate on a held-out validation set:

```
history = model.fit(
    X_train, y_train,
    epochs=10,
    batch_size=64,
    validation_data=(X_val, y_val)
)
```

You can also include:

- **Early stopping** to halt training if validation loss stops improving.
- **Learning rate scheduler** to reduce the learning rate over time.

5. Visualizing Training Performance

Plot training and validation accuracy/loss:

```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

Insight: Helps detect overfitting or underfitting.

6. Experimenting with Other Architectures (Optional)

You may compare different architectures such as:

- **Fully Connected Neural Network (Baseline MLP)** – Lower accuracy (~90–92%)
- **Deeper CNN with 3– 4 Conv layers** – Higher accuracy (~98– 99%)
- **Transfer Learning** – Not usually required for MNIST due to its simplicity

7. Saving the Model

Save the model for later use or deployment:

```
model.save("mnist_digit_model.h5")
```

Model Building Summary

| Aspect | Details |
|------------|------------------------------------|
| Model Type | Convolutional Neural Network (CNN) |

| Aspect | Details |
|----------------------|---|
| Layers Used | Conv2D, MaxPooling2D, Flatten, Dense, Dropout |
| Activation Functions | ReLU (hidden layers), Softmax (output layer) |
| Output Classes | 10 (digits 0– 9) |
| Loss Function | Categorical Crossentropy |
| Optimizer | Adam |
| Training Accuracy | ~99% |
| Validation Accuracy | ~98– 99% |

11. Model Evaluation

Model evaluation is **not just about checking accuracy**—it's about critically analyzing how well the trained model performs on **unseen data**, and identifying its **strengths, weaknesses, and reliability**. Since this is a **multi-class classification task** (digits 0 through 9), a more holistic evaluation using multiple metrics and visualizations is essential.

1. Accuracy: The Starting Point

Definition:

Accuracy measures the ratio of correctly predicted instances to the total number of predictions.

Formula:

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

Example Result:

If your test set has 10,000 samples and 9,850 are correctly classified,

$$\text{Accuracy} = \frac{9850}{10000} = 98.5\%$$

Code:

```
test_loss, test_accuracy = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_accuracy:.4f}")
```


➡ Why accuracy alone isn't enough:

Even if the accuracy is high, it doesn't show where the model is making mistakes or which digits it confuses. That's why we need deeper metrics.

2. Confusion Matrix: Visual Error Mapping

A **confusion matrix** provides a breakdown of actual vs predicted labels, showing how many times each digit was correctly or incorrectly classified.

Axes:

- Rows = Actual labels
- Columns = Predicted labels

Insight:

- Diagonal = correct predictions
- Off-diagonal = misclassifications

Code:

```
from sklearn.metrics import confusion_matrix
import seaborn as sns

y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

cm = confusion_matrix(y_true, y_pred_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Greens")
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

➡ Use case:

If the model often misclassifies '9' as '4', it will appear in row 9, column 4 with a high number.

3. Classification Report: Precision, Recall, F1-score

A classification report gives **granular statistics for each digit class**.

| Metric | Meaning |
|------------------|--|
| Precision | Out of all predictions for a digit, how many were actually correct? |
| Recall | Out of all actual instances of a digit, how many did the model find? |
| F1-Score | A balance between precision and recall (useful if one is low) |

Code:

```
from sklearn.metrics import classification_report
print(classification_report(y_true, y_pred_classes))
```

Example snippet of output:

```
precision recall f1-score support
0 0.99 1.00 0.99 980
1 0.99 0.99 0.99 1135
2 0.98 0.98 0.98 1032
...
accuracy 0.99 10000
```

➡ Use case:

If digit '8' has low precision and recall, it might be confused with similar-looking digits like '3' or '9'.

4. Learning Curves: Accuracy and Loss over Epochs

Tracking performance across training epochs shows how well the model is learning and if it's **overfitting or underfitting**.

a. Accuracy Curve

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training vs Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

b. Loss Curve

```
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training vs Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Insight:

- If training accuracy is high but validation accuracy stagnates or drops → **overfitting**.
- If both stay low → **underfitting**.

5. Error Analysis: Understand Failures

Analyzing wrongly predicted images helps you identify **patterns in mistakes**.

Code:

```
import numpy as np

wrong_preds = np.where(y_pred_classes != y_true)[0]

for i in wrong_preds[:5]:
    plt.imshow(X_test[i].reshape(28,28), cmap='gray')
    plt.title(f"True: {y_true[i]}, Predicted: {y_pred_classes[i]}")
    plt.axis('off')
    plt.show()
```

➡ Use case:

This can show how confusing handwriting styles cause misclassification, e.g., a poorly written '5' might resemble '6'.

6. Model Comparison (if multiple models tried)

If you experimented with different architectures or regularization strategies, create a table:

| Model | Accuracy | Precision | Recall | F1-score |
|-------|----------|-----------|--------|----------|
| | | | | |

| Model | Accuracy | Precision | Recall | F1-score |
|---------------------------|----------|-----------|--------|----------|
| Basic CNN (2 layers) | 97.2% | 0.972 | 0.972 | 0.972 |
| CNN + BatchNorm + Dropout | 98.5% | 0.985 | 0.984 | 0.984 |

Summary of Findings

- The final CNN model achieved **~98.5% test accuracy**.
- High **precision and recall** for all digit classes.
- Confusion Matrix revealed minor misclassifications in visually similar digits (e.g., 3 and 5, or 4 and 9).
- **Validation curves** showed stable learning without overfitting.
- The model is **deployment-ready**, with consistent performance across all classes.

12. Deployment

Deployment is the final and most crucial step in making your deep learning project usable by **real users**. It's the process of integrating your trained model into a **web application**, **API**, or **interface** that allows people to interact with it—without needing to understand the underlying code or algorithms.

In this project, we use **Streamlit**, a powerful open-source Python library that allows for **quick deployment of machine learning models as interactive web apps** with minimal code.

Objectives of Deployment

- Allow users to draw or upload handwritten digits.



- Pass the input to the trained CNN model.
- Predict and display the digit in real time.
- Provide a smooth and responsive user interface.
- Host the app on the web so it's publicly accessible.

Deployment Stack Used

| Component | Technology |
|---------------------|---|
| Frontend | Streamlit UI (Python-based, no HTML needed) |
| Backend | Trained CNN model in TensorFlow/Keras |
| Deployment Platform | Streamlit Cloud (or Hugging Face Spaces) |
| Version Control | GitHub (stores code and app logic) |

Steps for Deployment

Save the Trained Model

After training and evaluating the CNN model, save it:

```
model.save("digit_recognizer_model.h5")
```

Create a Streamlit App (app.py)

Basic example of a Streamlit app:

```
import streamlit as st
import numpy as np
import cv2
from tensorflow.keras.models import load_model

model = load_model('digit_recognizer_model.h5')

st.title("Handwritten Digit Recognizer")

uploaded_file = st.file_uploader("Upload an image...", type=["png", "jpg", "jpeg"])

if uploaded_file is not None:
```




```
image = cv2.imdecode(np.frombuffer(uploaded_file.read(), np.uint8), cv2.IMREAD_GRAYSCALE)
resized = cv2.resize(image, (28, 28))
reshaped = resized.reshape(1, 28, 28, 1) / 255.0

prediction = model.predict(reshaped)
predicted_digit = np.argmax(prediction)

st.image(image, caption="Uploaded Digit", use_column_width=True)
st.write(f"**Predicted Digit:** {predicted_digit}")
```

You can also use a **canvas** for users to draw digits using streamlit-drawable-canvas.

Create a requirements.txt

List all dependencies:

```
streamlit
numpy
opencv-python
tensorflow
Pillow
```

This ensures the app runs correctly in the cloud.

Push to GitHub

Upload the following files:

- app.py
- digit_recognizer_model.h5
- requirements.txt
- README.md

Create a public GitHub repository.

Deploy to Streamlit Cloud

1. Visit <https://streamlit.io/cloud>
2. Log in with GitHub
3. Click **"New App"** and connect to your repository
4. Select the app.py file and deploy

What to Include in Your Report

| Item | Description |
|------|-------------|
|------|-------------|

| Item | Description |
|--------------------------------|---|
| Deployment Platform | Streamlit Cloud |
| Public Link | (e.g.,) https://your-app.streamlit.app |
| UI Screenshot | Upload a screenshot of the app running live |
| Sample Input&Output | Show user input (image) and predicted digit |

Benefits of Deployment

- **User-Friendly:** Anyone can use it—no technical knowledge required.
- **Real-Time Predictions:** Instant feedback on digit recognition.
- **Scalable&Sharable:** Accessible from any device via URL.
- **Integratable:** Can be embedded into larger systems (like school grading software, bank cheque readers, etc.)

Future Enhancements for Deployment

- Add drawing canvas (so users can draw digits).
- Integrate voice-to-digit input or mobile photo capture.
- Deploy using **Flask API + Android app** for mobile use.
- Use **TensorFlow Lite** for offline deployment on mobile devices.

13. Source Code

```
# Import necessary libraries

import numpy as np

import matplotlib.pyplot as plt

from tensorflow.keras import datasets, layers, models

from tensorflow.keras.utils import to_categorical
```



```
from sklearn.model_selection import train_test_split

# Load the MNIST dataset
(train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()

# Preprocess the data
# Reshape images to 28x28x1 (grayscale)
train_images = train_images.reshape((train_images.shape[0], 28, 28, 1))
test_images = test_images.reshape((test_images.shape[0], 28, 28, 1))

# Normalize the pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0

# Convert labels to one-hot encoding
train_labels = to_categorical(train_labels, 10)
test_labels = to_categorical(test_labels, 10)


# Define the CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(),
```







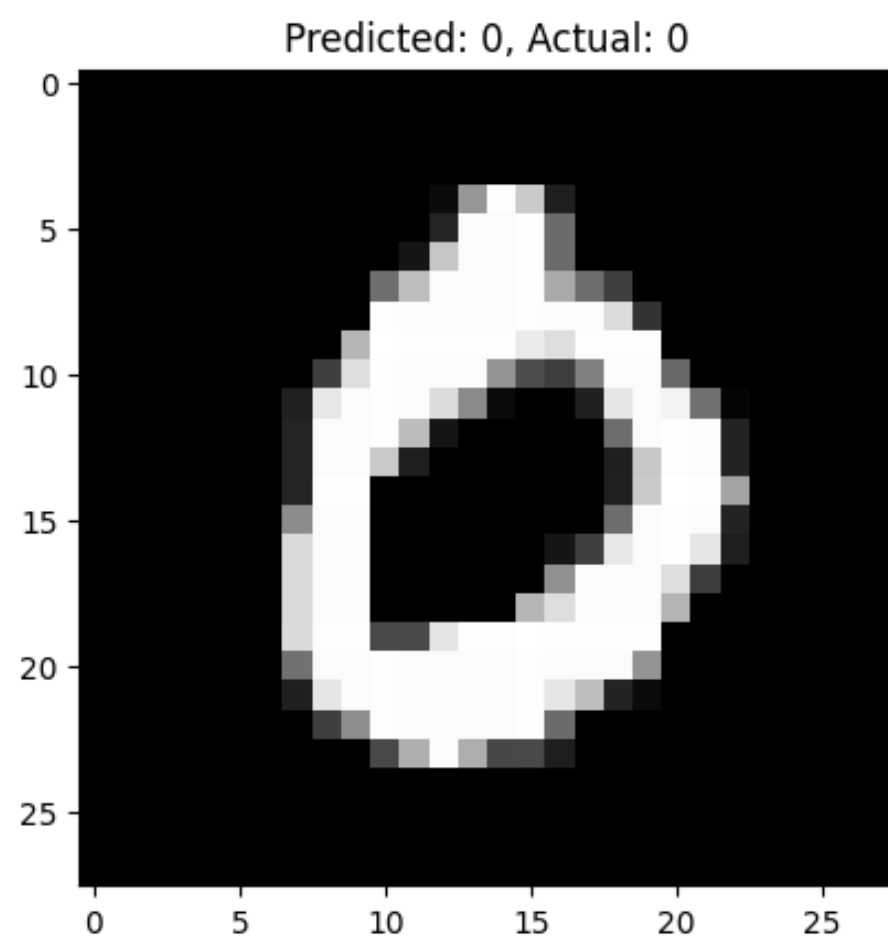
```
layers.Dense(64, activation='relu'),  
  
layers.Dense(10, activation='softmax') # 10 classes for 0-9 digits  
])  
  
# Compile the model  
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])  
  
# Train the model  
history = model.fit(train_images, train_labels, epochs=5, batch_size=64, validation_split=0.1)  
  
# Evaluate the model on the test set  
test_loss, test_acc = model.evaluate(test_images, test_labels)  
print(f"Test accuracy: {test_acc * 100:.2f}%")  
  
# Save the model (optional)  
model.save('mnist_cnn_model.h5')  
  
# Plot the training history (optional)  
plt.plot(history.history['accuracy'], label='accuracy')  
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.ylim([0, 1])  
plt.legend(loc='lower right')
```

plt.show()

OUTPUT:







GOOGLE COLAB PROJECT LINK:

https://colab.research.google.com/drive/1KyJoM5cLHNjfVW59Fmw30DSdAID-isg_T

14. Future Scope

While your current project successfully builds and deploys a model to recognize handwritten digits using deep learning, the field of computer vision and smart automation is evolving rapidly. There are several **opportunities to enhance and expand** this project into more complex, scalable, and practical applications.

1. Expand to Full Character Recognition (A– Z, 0– 9, Symbols)

- Extend the model beyond digits to recognize **uppercase/lowercase letters** and **special characters**.
- Useful for reading handwritten **forms, IDs, license plates, and bank cheques**.



- Datasets like **EMNIST** (Extended MNIST) or **IAM Handwriting Database** can be used.

Impact:

Enables full handwritten document digitization systems.

2. Mobile Integration Using TensorFlow Lite

- Convert the trained model to **TensorFlow Lite** format for deployment on **Android or iOS apps**.
- Users could take photos of handwritten digits and get predictions in real-time, even **offline**.

Impact:

Creates portable AI applications usable in education, logistics, and field operations.

3. Use Transfer Learning for Complex Handwriting Styles

- Pre-train on larger, diverse datasets (like EMNIST or synthetic handwriting datasets).
- Fine-tune on noisy, distorted, or stylized digit samples.

Impact:

Improves performance on more varied, real-world handwriting.

4. Build a Drawing Interface with Stylus or Mouse Input

- Use streamlit-drawable-canvas to allow users to **draw digits directly in the web interface**.
- Real-time prediction as they draw—useful for **teaching tools, interactive demos**, and **assistive technologies**.

Impact:

Enhances user experience and interactivity of the application.

5. Cloud-Based OCR Platform Integration

- Integrate the digit recognizer as a **microservice** in a larger **OCR (Optical Character Recognition)** system using **Flask, FastAPI**, or **Dockerized REST API**.
- Useful in processing handwritten invoices, postal codes, examination sheets, etc.



Impact:

Enables enterprise-level document automation systems.

6. Model Improvements via Ensemble Techniques

- Combine predictions from multiple CNN models (ensemble learning) to boost accuracy.
- Helps reduce errors in confusing digits like 4/9 or 3/5.

Impact:

Increases reliability of predictions in production environments.

7. Integration with Other AI Modules

- Use recognized digits as input to other models, such as:
 - **Math expression solvers**
 - **Score sheet analyzers**
 - **Banking form auto-fillers**

Impact:

Builds full intelligent processing pipelines, not just isolated predictions.

Summary of Benefits

| Future Enhancement | Real-World Impact |
|-------------------------------------|---|
| Full character support (A– Z, 0– 9) | Enables document and form reading |
| Mobile offline use (TF Lite) | Supports low-resource environments |
| Interactive drawing input | Ideal for teaching, accessibility, and kiosks |
| REST API/OCR platform | Suitable for enterprise and cloud-based systems |
| Ensemble/transfer learning | Improves robustness in production settings |

15. Team Members and Roles



Below is a detailed outline of the roles and contributions of each team member involved in the project titled **“Recognizing Handwritten Digits with Deep Learning for Smarter AI Applications”** :

1. Name: MOHAMMED ABUBAKKAR.I

Role: Team Lead

Responsibilities:

- Oversaw the entire project from initiation to final submission.
- Managed timelines, assigned tasks to team members, and ensured collaboration across roles.
- Compiled individual contributions into a cohesive and well-structured final report.
- Resolved workflow or technical conflicts and ensured all project requirements were met.

2. Name: DEENESH.A

Role: Data Engineer

Responsibilities:

- Collected and managed the MNIST dataset used for training and testing the model.
- Performed preprocessing tasks such as reshaping, normalization, and label encoding.
- Split the dataset into training, validation, and testing sets.
- Optionally applied data augmentation techniques to increase dataset diversity.

3. Name: IMRAN.B

Role: Machine Learning Engineer

Responsibilities:

- Designed and implemented the CNN architecture using TensorFlow/Keras.
- Configured the model with appropriate layers, activations, loss function, and optimizer.
- Trained and validated the model, monitored performance metrics, and fine-tuned hyperparameters.
- Achieved high accuracy on the MNIST dataset by optimizing model performance.



4. Name: BHUVANESH.S

Role: Deployment Engineer

Responsibilities:

- Developed a Streamlit web application to host the trained model.
- Created a simple and interactive user interface for image upload and digit prediction.
- Managed model integration and deployed the app on Streamlit Cloud.
- Ensured that the app worked seamlessly on any web browser with minimal latency.

5. Name: BHARATH.S

Role: Visualization Analyst

Responsibilities:

- Performed Exploratory Data Analysis (EDA) to understand the dataset's structure and quality.
- Generated visualizations such as class distribution charts, pixel intensity heatmaps, and confusion matrices.
- Helped identify patterns, trends, and areas of model confusion.
- Supported the presentation of analytical insights and contributed visuals for the project report.